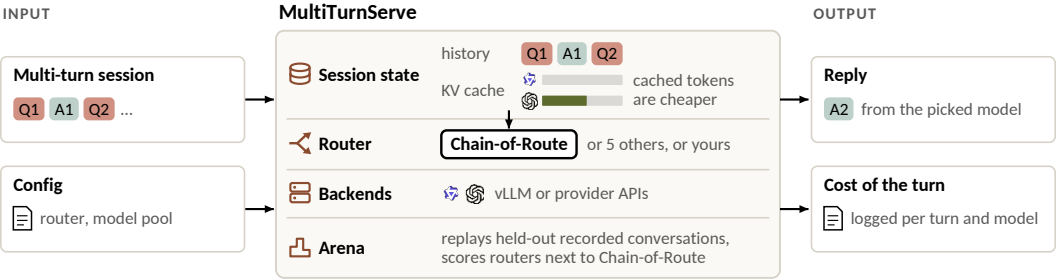
INPUT
MultiTurnServe
OUTPUT
Multi-turn session
Q1 A1 Q2 ...
Config
router, model pool
Session state
history
Q1 A1 Q2
KV cache
cached tokens are cheaper
Router
Chain-of-Route
or 5 others, or yours
Backends
vLLM or provider APIs
Arena
replays held-out recorded conversations, scores routers next to Chain-of-Route
Reply
A2 from the picked model
Cost of the turn
logged per turn and model