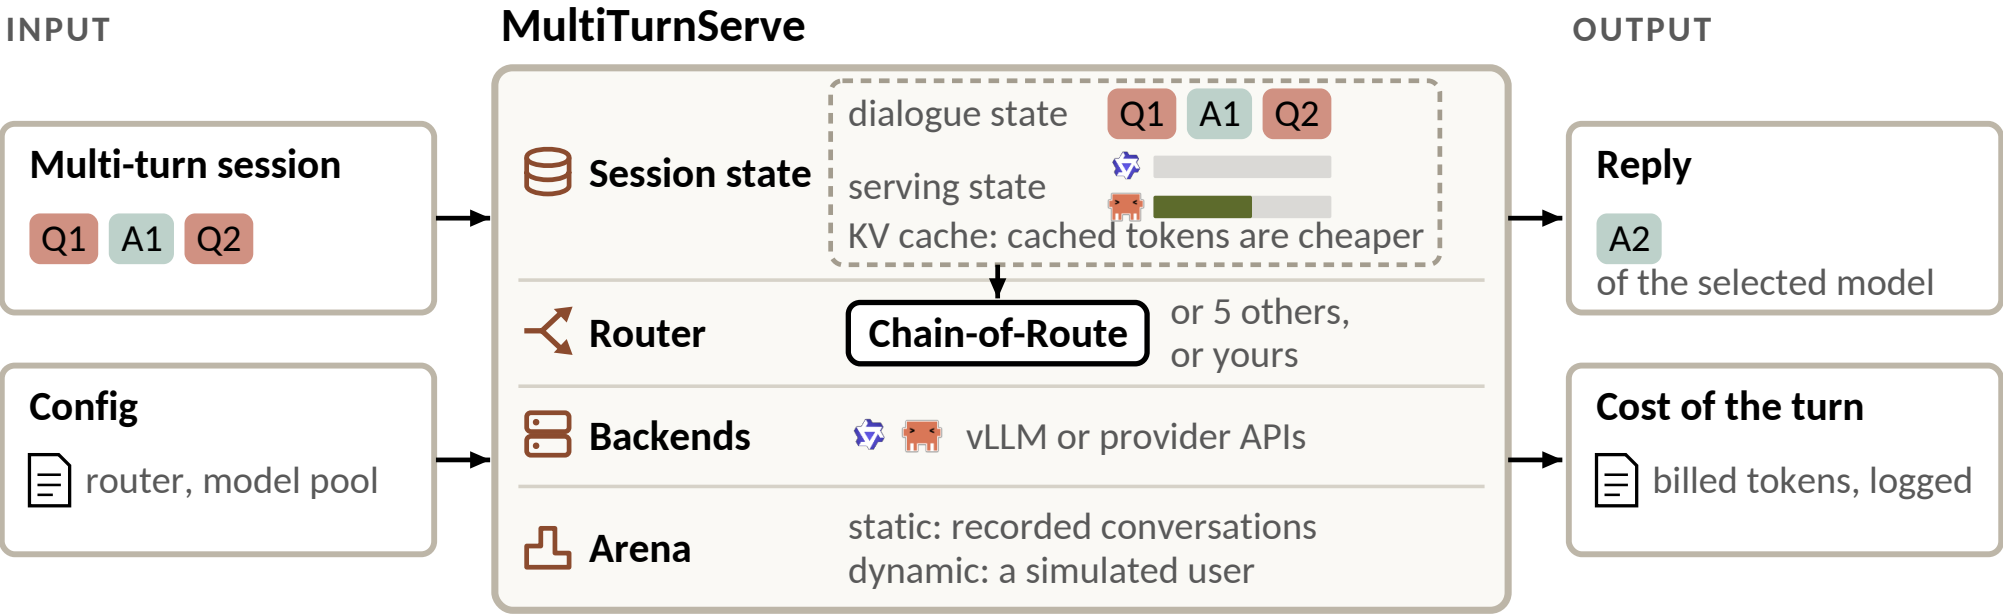
INPUT
MultiTurnServe
OUTPUT
Multi-turn session
Q1
A1
Q2
Config
router, model pool
Session state
dialogue state
Q1
A1
Q2
serving state
KV cache: cached tokens are cheaper
Router
Chain-of-Route
or 5 others,
or yours
Backends
vLLM or provider APIs
Arena
static: recorded conversations
dynamic: a simulated user
Reply
A2
of the selected model
Cost of the turn
billed tokens, logged